

## **Module 1**

### **Introduction and overview**

- Review of basic computer architecture, Technology trend and design goals, Performance metrics and performance enhancement techniques, Power dissipation in processors, power metrics, and low-power design techniques.

### **Design goals for high performance**

- Function requirements: process data according to instructions
  1. Tests and verifications are required at different stages since it is not possible to modify the processor once fabricated.
- Reliability: should continue to perform correctly
  1. Important for all modern computers and more important for mission-critical application
  2. Reliable and fault-tolerant computing is an important area of study, which considers various approaches to improve fault tolerance and reliability.
- Cost: a very important factor
  1. Embedded consumer products are particularly highly sensitive to cost
- Performance: is a basic requirement
  1. A computing system needs to provide the desired performance.
- Power consumption: is a very important requirement nowadays
  1. A system needs to consume less power for many reasons: battery life, cost of electricity, thermal problem, and cooling cost.
- **Performance and power consumption:** The evolution of computer architecture is driven to improve these two goals.

### **Iron law**

$$\text{Performance} = \frac{1}{\text{Execution Time}}$$

$$\text{Execution time} = IC \times CPI \times T$$

- A decrease in one may lead to an increase in the other two

### Speed up of a processor

$$\text{Speed up} = \frac{\text{Time original}}{T_{\text{enhanced}}} = \frac{1}{[(1-E)+E/S]}$$

### Amdahl's law

Speedup via parallelism is limited by that component of the application that cannot be enhanced; in a program with parallel processing, a relatively few instructions that have to be performed in sequence will have a limiting factor on program speedup such that adding more processors may not make the program run faster.

### Gustafson's law

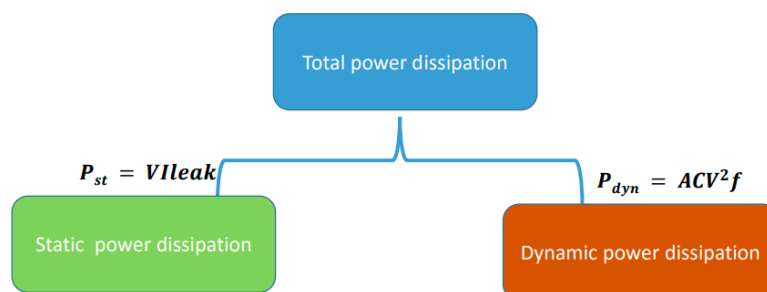
Completing more tasks in a given amount of time using more processors. Run each task on one processor, but do these task together ❤️

### MIPS as a performance metric

A million instructions per second (MIPS) is an approximate measure of a computer's raw processing power.

### Power dissipation in processor

Power metrics (Dynamic and static power)



Lower operating voltage  $V$

$A$ : switch factor,  $C$ : capacitance,  $V$ =voltage,  $f$ =clock frequency

### Low power design techniques

- Efficient cache and memory hierarchy design
- Power gating: shutting down the unused components(enable signals)
- Clock gating: to reduce unnecessary switching
- Reducing the data movement, number of memory access, and register transfer

## **Module 2: Instruction set architecture design**

### **Introduction to ISA**


A set of following specifications that a programmer must know to write a correct and efficient program for a specific machine.

- Instruction format
- Length of instruction and size of the field
- Word size: 16-bit, 32-bit or 64-bit
- Set of all operations: opcodes/machine language
- Register file specification: size, width, and its usage of registers in CPU
- Memory address space and addressability: no of addressable locations & bits per location
- Addressing modes: ways of specifying and accessing operands(s): indicate how and address is determined
- Operand locations: all in registers, register and memory or all in memory

*ARMv8 ISA: A design example named LEGv8*

### **Instructions classified based on functionality**

- **Register(R) Type:** fetching the data from two register and perform an operation, store the result to the new register



- All data values are located in the register
  - **Rn & Rm:** specify the first and the second source register
  - **Rd:** specifies the destination register
  - **Shamt** stands for shift amount: specifies the number of bit positions to be shifted
  - **Opcode:** type of instruction
- **Data Transfer(D) Type(Load/Store):** load: data will be fetched from one location to register/ store: data be fetched then store in another location with the help of registers.

opcode	address	op2	Rn	Rt
11 bits	9 bits	2 bits	5 bits	5 bits

- **Rn:** specify the base registers
- **Address:** constant offset from contents of base register
- **Rt:** specifies the destination register (load) or as the **source**(store) register number(contain data to be stored into **Rn +address**)
  - ★ ALU will calculate the address(address +[Rn])
  - ★ Read the Data at memory location(address + [Rn])

- **Immediate (I) Type:** add an immediate number with **Rn** and save it into **Rd**

opcode	immediate	Rn	Rd
10 bits	12 bits	5 bits	5 bits

- **Rn:** Source register
- **Rd:** destination register

- **Conditional instructions:** branch to a labeled instruction if a condition is true, otherwise just continue
  1. **CBZ register, L1** : if(register ==0) branch to instruction labelled L1
  2. **CBNZ register, L1** : if(register !=0) branch to instruction labelled L1
  3. **B L1**: branch unconditionally to instruction labelled L1

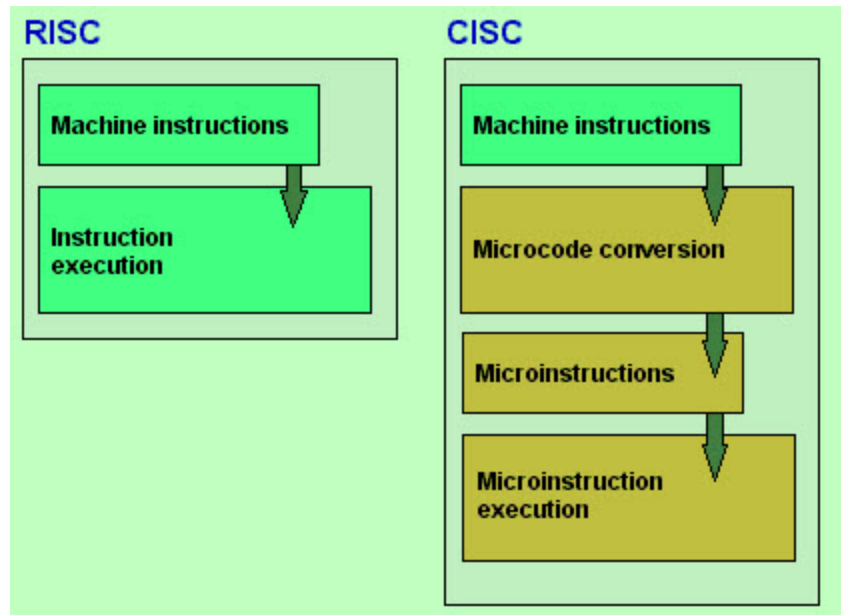
Relative advantages of RISC and CISC instruction set

### CISC

- Many complex instructions in the instruction set
- Many formats, & several addressing modes to support complex instructions
- Instruction length varies according to the addressing mode

### RISC

- Fewer simple instructions in the instruction set
- Fewer instruction formats, & a few addressing modes
- Fixed instruction length ) simpler implementation

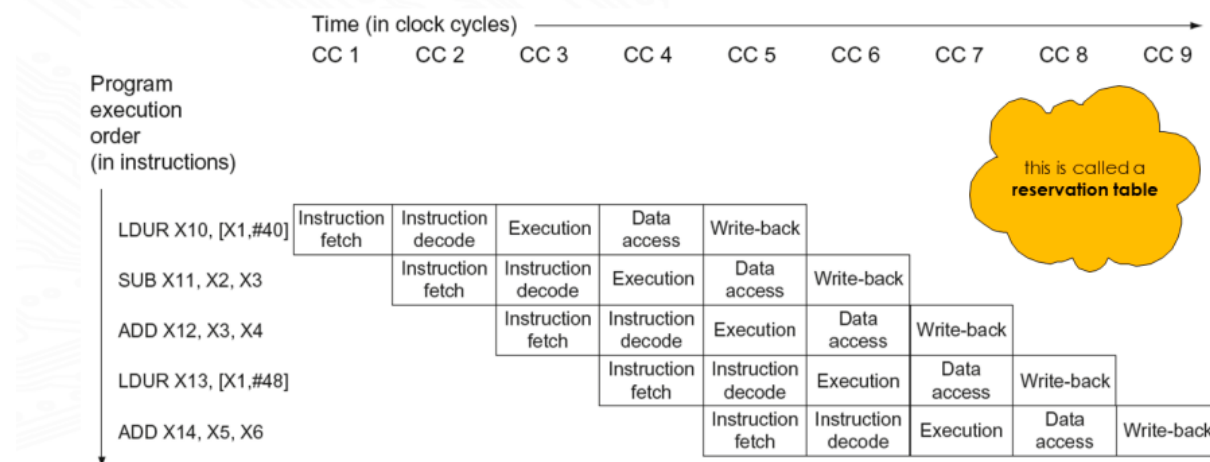


## Module 4: Pipelining and Instruction Level Parallelism

Visualization of pipeline architecture

**Fetch Decode Execute Memory access Write-back**

Example of ideal pipelining



Each row of **reservation table** corresponds to a **pipeline stage** and each column represents a clock cycle.

Challenges in the realization of pipeline architecture

- **Ideal pipeline**
  1. Identical task of all instructions
  2. Uniform decomposition of task
  3. Independent Computations

- **Real pipeline**
  1. Except for load instruction, all other instructions do not need all stages
  2. Not all pipeline stages involve, using the same time to complete their respective sub-task
  3. The execution of one instruction may depend on one of the preceding instructions
- **Complications**
  1. Datapath
    - a. Many instructions in flight
  2. Control
    - a. Must correspond to multiple instructions
  3. Instructions may have
    - a. Data and control flow dependencies
    - b. One may have to stall and wait for another

#### Pipeline Hazards

- **Stalling the pipeline will result in a drop in efficiency**
- **Structural hazard:** if two instruction needs the same piece of hardware resource(eg memory) at the same time, a stall will happen, one instruction needs to wait for at least one clock cycle.
  - **Solution:** get additional hardware elements
- **Data hazard:** either the source or destination register of instruction is not available at the time expected by the pipeline, causing a stall
  - **Solution:** know the dependencies between data in the program and eliminate them
- **Control/Instruction hazard:** conditional jumps or jumps or subroutine calls can stall a pipeline because of the delay in the availability of instruction.
  - **Solution:** Need to predict the location of branches to avoid stalling

#### Data-dependence and Handling data-dependence

- Types of data dependence
  - **Flow dependence > RAW:** read after write
    - Detect and wait until the value is available in register file
    - Detect and bypass data to dependent instruction
    - Detect and eliminate the dependence at the software level
    - Do something else
  - **Output dependence > WAW:** write after write
  - **Anti dependence > WAR:** write after read

### Out of order execution and dynamic scheduling

- In-order instruction execution
  - Instructions are fetched, executed & completed in compiler-generated order
  - One stall, they all stall
  - Instructions are statically scheduled
- Out-of-order instruction execution
  - Instructions are fetched in compiler-generated order, but they may be executed in some other order
  - Independent instructions behind a stalled instruction may pass it
  - Instructions are dynamically scheduled during run time
- Dynamic scheduling
  - If there are no structural or data hazards(hardware or register conflicts), the instruction can bypass and execute its instruction first.

### Instruction reordering and Register renaming

- Reorder instructions to maximise efficiency
- Register rename: need more registers

### Loop unrolling

Change loops into a longer code sequence

- Greater demand for registers
- Minimize structural hazards and data hazards

### Control hazards handling

When branching, the pipeline will have to stall as it changes the sequence of instructions to be executed

- Instruction fetched previously by the instruction fetch unit needs to be discarded and replaced with the new instruction resulting in stall cycles.
- Time lost due to branching is referred as **branch penalty**
  - **Solution:** stall the pipeline immediately after we fetch a branch instruction; **3 stalls** when Branch Target Address is updated in the Memory stage.

### Static and dynamic branch prediction

- **Branch prediction:** assume the branch does not happen and continue executing the instructions in the pipeline. If the branch does occur, the pipeline must be flushed of the speculative instructions.
  - Static branch prediction: when behavior is highly predictable

- Branch always taken: make sense for if pipeline location is known before the branch outcome > the pipeline will execute the instructions in the branch while waiting for the outcome
  - Branch always not taken: speculation
- Dynamic branch prediction: prediction decision may change depending on the execution history

### Instruction level parallelism

Parallel execution of instructions requires 3 major task

- Check for data dependence between instructions to see if the instructions can be grouped together for parallel execution
- Assigning instructions to different functional units in the processor
- Determine when the instruction is to be executed
- Methods to extract more parallelism
  - Instruction reordering and out of order execution: change the order of execution of instruction if it does not violate the data dependence
  - Speculative execution with dynamic scheduling: to execute an instruction without exactly knowing if that needs to be executed: ahead of branch outcome (do first then check)
  - Loop unrolling

### Concepts of superscalar processing

(down to each core, maximize the use of every functional unit in the core)

- Multiple Instruction fetch in parallel
- Multiple instructions decode in parallel
- Requirements:
  - Multiple functional units in **a CPU** are used
  - More than 1 independent instruction needs to be available to be executed
  - Use specialised hardware to issue independent instructions that can be executed simultaneously

### VLIW architecture (advantages and disadvantages)

- Packs multiple independent operations into one instruction.
- VLIW processor needs a compiler to break the program instructions down into basic operations that can be performed by the processor in parallel
- These operations are put into a very long instruction word which the processor gets executed in appropriate FUs.
- There must be enough parallelism in the program code to fill the available slots, to reduce the CPI significantly
- **Advantage:**



- The compiler will take care of all the scheduling and plans for execution
- **Disadvantage:**
  - Compatibility across implementations is a major problem.
    - Cannot run properly with different number of FU or FUs with different latencies
    - Unscheduled events may stall the entire processor

## Module 5: Memory Systems

### Cache design

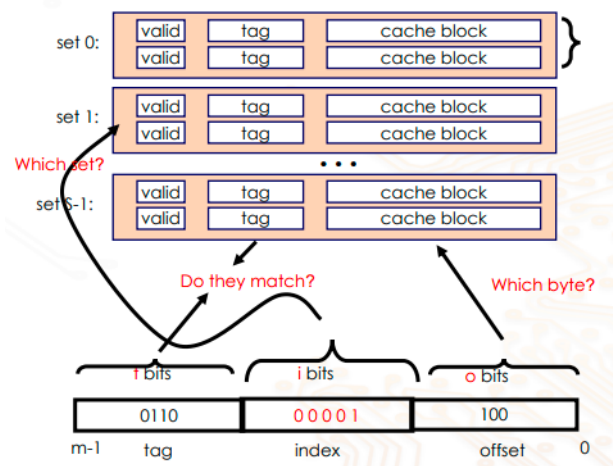
- Level 1 cache: small, resides in the processor
- Level 2 cache: bigger than L1, resides between CPU and main memory, is slower than L1 cache.
- Needs to convert the main memory address into a cache location.

Consider 64-bit Address: eg: ARMv8

Portion	Length	Purpose
Offset	$o = \log_2(\text{block size})$	Select word within block/line
Index	$i = \log_2(\text{number of sets})$	Select the set
Tag	$t = (\text{address size}) - o - i$	ID block/line within set

- Placement policy: **cache organization schemes**
  - Direct-mapped cache
    - Simplest and fastest address mapping
    - A memory block is mapped to a specific cache line only

- Cache consists of an array of fixed size frames called cache line
- Each frame holds a block of main memory
- Direct mapping: each memory block is mapped to a specific cache line
- A set of memory blocks with the same cache index are mapped to the same cache line.
- Cache size (roughly):  $C = BS \times S$  data bytes
  - Index in the address will be used to locate the **set** in the cache
  - Tag in the address will be used to check if the cache set is valid(cache hit) (if tag are equal)
  - Offset in the address be used to find the block in that cache line
- Fully associative cache
  - A memory block can be placed in any of the cache lines.
  - Flexible but expensive
    - Need to search all the cache
    - Search the requested block tag with every tag in the cache to find the desired block
    - Main memory only need to be partitioned as tag and offset
    - Expensive in terms of speed and complexity
- Set-associative
  - Combines the strategy of fully associative and direct-mapped caches
  - Similar to direct mapping but more lines per set



### Cache replacement and write policies

**Replacement:** Choosing which cache line to replace with new memory data

- Several policies:
  - FIFO
  - LRU(least recently used): takes care of temporal locality
    - Needs to keep a history of access and hence slows the system down
  - NMRU(Not most recently used)
  - Pseudo-random

### **Write policies**

- Memory hierarchy
  - Two or more copies on the same block
  - Main memory or disk
  - Cache
- What to do on a write
  - All copies need to be changed
  - The cache must propagate to all levels of the memory hierarchy
- **Write Through:**
  - Write propagates directly through the hierarchy (L1, L2, memory, disk)
  - High bandwidth requirement
  - Memory becomes slow when its size increases
  - Popular in the system only for writing to the L2 (beyond L2, use write back)
  - Defeats the purpose of having a cache which is to reduce access to the main memory
- **Write Back:**
  - Maintain state of each line in a cache
    - Invalid = not present in the cache
    - Clean = present, but not written
    - Dirty = present and written
  - Store state in tag array, next to address tag
    - Mark dirty bit on a write
  - On eviction/when the cache line is going to be cast out
    - If the dirty bit is set, write back the dirty bit to the next level
- Complications in write policies
  - More stale copies in the lower level of the hierarchy
  - Must always check higher for dirty copies before accessing copy in a lower level, because dirty copies are being updated in the higher level before lower-level cast out.

- Cache coherence problem in multiprocessors
- I/O devices that use direct memory access have to check the cache first before reading the main memory
- Instruction cache( better to have it apart from data cache)
  - Stores instructions only.
  - Tends to have high locality and thus a high hit rate

### Cache performance analysis

#### Cache

- Cache hit: life is made easy and improves performance
- Cache miss: fetch from next level

#### What is a performance impact?

- Miss penalty leads to a reduction in performance
  - Detect miss: 1 or more cycles
  - Find victim(line to be replaced)
    - Write back if the victim is dirty
  - Request line from next level: several cycles
  - Transfer line from next level
  - Fill line into data array, update tag array
  - Resume execution
- A cache miss is determined by
  - Temporal locality
  - Spatial locality
  - Cache organization
- Reasons for a cache miss
  - Compulsory miss
    - First-ever reference to a given block of memory
  - Capacity miss
    - Working set exceeds cache capacity
    - Useful blocks (with future references) displaced
  - Conflict miss
    - Placement restrictions (not fully assoc). Cause useful blocks to be displaced
- Cache miss effects
  - More block is better, since greater capacity, fewer conflicts
  - Associativity
    - Higher associativity reduces conflicts
    - Very little benefits beyond 8 way set associative

- Block size
  - Larger blocks exploit spatial locality
  - Usually, the miss rate improves until 64B-256B
  - 512B or more miss rates get worse
    - Fewer placement choices: more conflict misses
    - High miss penalty

$$\begin{aligned} \text{CPU time} &= \text{IC} \times \text{CPI} \times \text{cycle time} \\ &= \text{IC} \times (\underbrace{\text{CPI}_{\text{ideal}} + \text{Memory-stall cycles}}_{\text{CPI}_{\text{stall}}}) \times \text{cycle time} \end{aligned}$$

- Memory-stall cycles = memory accesses/program  $\times$  miss rate  $\times$  miss penalty
- A processor with a  $\text{CPI}_{\text{ideal}}$  of 2, a 100 cycle miss penalty,
- 36% load/store instr's, and 2% IM and 4% DM miss rates

Memory-stall cycles = Memory stall cycles in IM + memory stall cycles in DM

Memory stall cycles in IM =  $2\%$  (miss rate of IM)  $\times$  100 (Miss penalty) = **2**  
(All instructions has to be fetched from IM)

Memory stall cycles in DM =  
 $36\%$  (LW/SW)  $\times$   $4\%$  (miss rate)  $\times$  100 (miss penalty) = **1.44**

So  $\text{CPI}_{\text{stalls}} = 2 + 2 + 1.44 = \mathbf{5.44}$  (more than twice the  $\text{CPI}_{\text{ideal}}$  !)

## **Module 6:GPU Architecture and CUDA Programming**

### Motivation for GPU programming

GPUs are designed to accelerate 3D graphics renderings on PC

GPUs are also used extensively for many mathematical and scientific purposes

### Evolution of GPU architectures

Before GPUs, graphics display is based on a simple framebuffer subsystem standard known as the VGA

The original design of the graphics processing unit

- Based on fixed-function pipelines
- Configurable but not programmable
- Inflexible

### Programming on GPU (CUDA)

Nvidia released CUDA(compute Unified Device Architecture)

- A software platform that gives direct access to GPU's internal processing units to support parallel programming
- General-purpose parallel computing programming model where a problem is first expressed in the form of multiple threads
  - And execute in parallel using CUDA cores in the GPU
  - Enable many complex computational problems to be solved in a very efficient way
    - First decomposed into subproblems each of which is allocated to a "thread block"
    - Each subproblem is then separated into finer pieces that can be solved cooperatively in parallel by using multiple threads within the thread block
- CUDA Programming Model
  - CUDA operates on a heterogeneous programming model that consists of a host and a device
    - Typically a CPU and its memory(host)
    - The device is GPU and its memory(device)
  - A program will start and be run by the host
    - Which then launches the parallel threads that are executed on the device

Basic internal operations of GPUs

## Device Code (kernel)

To indicate the code that is to run on the device

- use the CUDA C keyword `__global__` declaration specifier

```
1  __global__ void hello_GPU(void) {  
2      printf("Hello from GPU!");  
3  }
```

Device code are launched as **Kernel** in CUDA

- which can be called from the host code as follows

```
4  int main(void) {  
5      hello_GPU<<<1,1>>>();  
6      printf("Hello from CPU!");  
7      return 0;  
8  }
```

- During compilation, the source file will be split into host and device components
- <<<x,y>>>: x means the number of thread blocks, y means the number of threads in every block, so xy amount of parallel threads currently concurrently
- **Synchronization between host:**
  - We need the host to finish the kernel execution using CUDA function
    - `cudaDeviceSynchronize()`
    - `cudaMalloc()` #allocate memory
    - `cudaMemcpy()` #sending data to device
    - `cudaMemcpyHostToDevice`
    - `cudaMemcpyDeviceToHost`
- `__syncthreads()` : synchronize between Threads
- Block vs Threads
  - Not ideal to have 1 thread per block or only one block with multiple threads

## Module 7: Data and Thread Level Parallelism


### Three levels of parallelism

- Instruction-Level parallelism
  - Multiple independent instructions are identified and grouped to be executed concurrently in different functional units in a single processor
  - Can reduce CPI to values less than 1
  - Eg. Superscalar and VLIW processor
- Data-level parallelism
  - The same operation is performed on multiple data values concurrently in multiple processing units. Can reduce the Instruction Count to enhance performance. Examples: vector processors and array processors
- Thread/Task Level parallelism
  - More than one independent thread/task is executed simultaneously.
  - Can reduce the total execution time of multiple tasks

### Data-level parallelism and SIMD processors

#### **Data level parallelism**

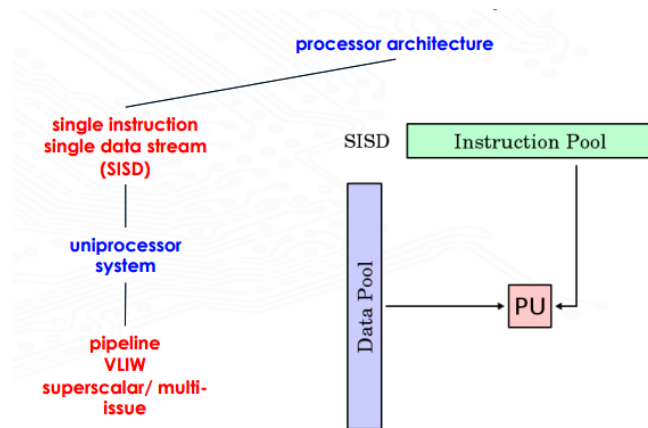
- Divide the parts of data between different tasks and perform the tasks in parallel
- **Key:** no dependencies between tasks that cause their results to be ordered
- Eg. Fig A
- Eg. Vector addition



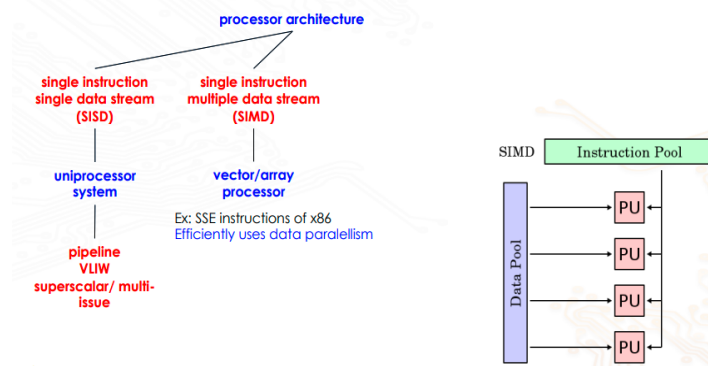
*Fig A*

- Computing models-Flynn's classification
  - Single instruction, single data stream(SISD)
    - Conventional sequential processor
    - Not suitable to realize data-level parallelism

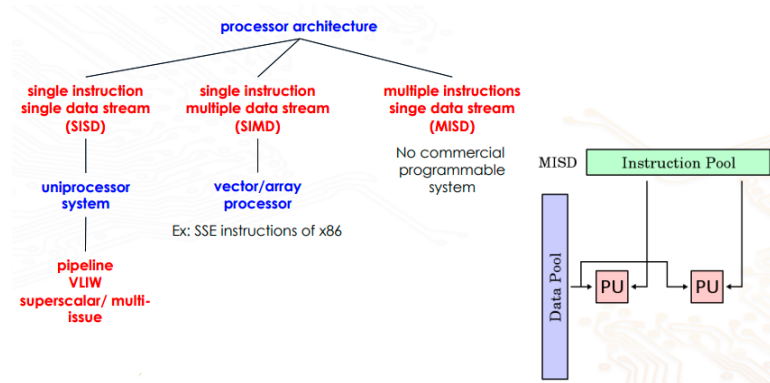




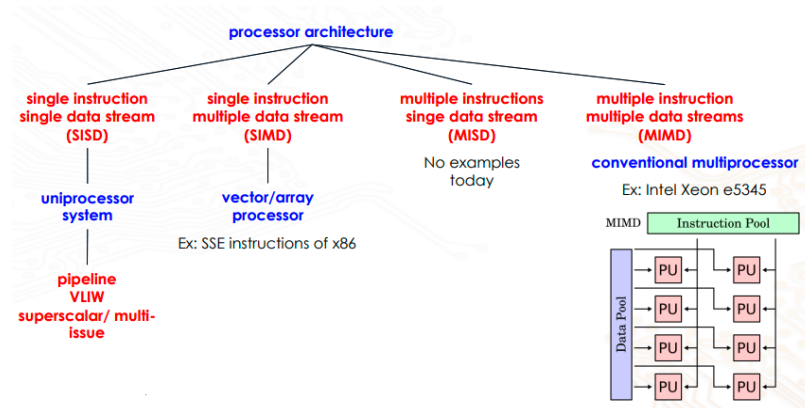
- Single instruction, multiple data stream(SIMD)
  - Efficiently utilize the data-level parallelism



- Multiple instructions, single data stream(MISD)
  - No commercial programmable system \*MISD structure is only of theoretical interest since no practical system has been constructed using this organization.



- Multiple instructions, multiple data streams (MIMD)
  - True multiprocessor system that issues multiple instructions
  - Can support data-level parallelism



- **Single instruction multiple data stream (SIMD)**
  - One instruction is issued every cycle, the same instruction is used to execute for different data elements by different processors
  - The central control unit issues instructions and controls the simultaneous execution of instructions.
  - Each processing unit has its own data memory and also could interact via shared memory.
  - All processing units work synchronously
  - Application: signal processing, matrix/scientific computation.
  - Examples: vectors and array processors
  
- **Advantages of SIMD over MIMD**
  - SIMD make use of data-level parallelism efficiently.
  - If there are multiple processing units, N- operations can be performed in one clock cycle =>high performance
  - Loop-overhead instructions(address increment and branch condition check) will also be reduced by a factor of N
    - **overhead** is any combination of excess or indirect computation time, memory, bandwidth, or other resources that are required to perform a specific task
  - Energy efficiency: energy used in instruction fetch will be reduced since there is only one instruction
  - The biggest advantage is that the programmer continues to think sequentially yet achieves parallel speedup by having parallel data operation
  
- **SIMD Processors**
  - Vector processor
    - Same operation at the same processing element (this PE will do + only)
    - Different operations at the same time(at time t, one PE can only handle one operation)
  - Array processor


- Same operation at the same time(a few PE will do the same operation at the same time)
- Different operations at the same processing element(every PE can do different operations, but cannot do them all at the same time)

### Thread-level parallelism and multi-core processors

Doing many tasks at the same time: running 2 programs concurrently

- Motivation for multicore
  - Need for performance enhancement
    - Need a faster and more capable system to support increasing transmission speed and communication bandwidth
    - Increase resolution of image/video data
    - Increase the use of video conferencing and surveillance.
    - New applications: Drug discovery with a high volume of DNA data processing, scientific computing, and weather forecasting
- Moore's law
  - The number of transistors in an IC chip doubles approximately every 1.5 years
- Traditional approaches for performance enhancement
  - Increase clock speed
    - Not really a good idea: might overheat the CPU, damage Processors, or Excessive Cooling that would take more time
  - Decrease clocks per instruction(CPI)
    - Instruction level parallelism
- The ILP wall


### The ILP Wall



- High design and verification time and cost
- Diminishing returns in more ILP
- Instruction-level parallelism is near its limit

- The Power Wall

## The Power Wall





$$P = \alpha \cdot C \cdot V_{dd}^2 \cdot f$$

$V_{dd}$  increases with  $f$

- Power consumption increases more rapidly with the increasing operating frequency
- When frequency increases, the power dissipated dynamically increases as well

- The Memory Wall

## The Memory Wall



- The rate of improvement in microprocessor performance far exceeds the rate of improvement in DRAM memory speed. Such a trend rendered the memory subsystem to become one of the most crucial system-level performance bottlenecks.

- The Brick Wall

- Can be overcome by having more cores/ multicore processor
  - Able to do more task simultaneously
  - Consumes lesser power due to lower clock frequency

- Advantage of multicore

- Overcome power wall
- Derive parallelism by thread-level parallelism
- **Homogenous multicore systems:** consist of identical cores, **less design and verification cost**
- **Heterogeneous multicore systems:** consist of different types of cores each optimized for different roles

Challenges:

- **A New Power Wall**
  - Moore's law, more transistor density, more heat dissipated per unit area
- **A New Memory Wall**
  - Due to a lack of **\*memory bandwidth**
    - *the rate at which data can be read from or stored into a semiconductor memory by a processor.*
  - Contention between cores over the memory bus